
PARALLEL SYSTEMS EXERCISE-2A

LABORATORY 2024-25 (15%)

Write a program in **OpenMP** that sorts a sequence of N integers $A[0\dots N-1]$ in parallel, according to the recursive sorting algorithm '**multisort**' listed on slides #49-50 of the OpenMP.ppt file .

[Consult among others and rely for your implementation on the code of the file **OpenMP_msrt.c** (indicative implementation of the mergesort sorting algorithm in OpenMP) posted on Eclass.]

Measure the performance of your implementation for large values of 'N' and for different numbers of threads, and compare them against each other. acceleration achieved in each case.

Deliverables: code, annotation/documentation, sample runs/results.

PARALLEL SYSTEMS EXERCISE-2B

LABORATORY 2024-25 (25%)

1. Write a program in **CUDA** that (given as input 'A', 'N'), finds for a two-dimensional array of integers $A(NxN)$, the average value of its elements (m), and its maximum ($amax$) element. Then check if its maximum element is greater than N -times its average value ($amax > N*m$), and accordingly:

- If this is true, create a new matrix $B(NxN)$ (which it prints to the screen at the end) where:

$$B_{ij} = (m - A_{ij}) / amax$$

For the above table B , it is also requested to calculate and print on the screen, the minimum element value of ($amin$).

- If this is not the case, create a new matrix $C(jxN)$ (which it will print on the screen at the end) where:

$$C_{ij} = (A_{ij} + A_{i(j+1)} + A_{i(j-1)}) / 3$$

Also consider that if $j+1=N$ then $A_{i(j+1)}=A_{i0}$, while if $j-1=-1$ then $A_{i(j-1)}=A_{i(j-1)}$.

2. Write a program in **CUDA** that, given a two-dimensional array, integers $A(NxN)$ calculates the covariance log of A as follows:

- For each column of table A , it calculates the average of the column's elements.
- From each element of a column of table A subtracts the average of the corresponding column calculated in the previous step.
- Multiples the matrix A of the previous step by its inverse. Since the resulting matrix is symmetric, it is sufficient to calculate either the upper or lower triangular part of the result.

Other Requirements: Measure the performance of your implementation for different values of 'N', different total number of threads (e.g. N, N^2 etc.), and different number of blocks and threads/block. Also consider the possibility of using shared memory in each subquery. Finally, where critical region protection / reduction is required, consider the possibility of using both *atomic commands* and the possibility of applying a *binary tree algorithm*.

Deliverables: Code, annotation/documentation, sample runs/results.

Submission Deadline: 12/1/2025 (pp. you are invited to work on only one of your choice – or if you wish, both – of the two parts of the Exercise – see the relevant Announcement posted on Eclass for more details). The Exercises you will submit will be examined orally at the end of the semester.